

COGS 205B: Computational Lab Skills for Cognitive Scientists I

Joachim Vandekerckhove

Spring 2022

Part 7: Metropolis-Hastings sampling

Metropolis-Hastings sampling

- A widely applicable algorithm for numerical integration is the **Metropolis-Hastings** (MH) algorithm

Metropolis-Hastings sampling

- A widely applicable algorithm for numerical integration is the **Metropolis-Hastings** (MH) algorithm
- In the MH algorithm, we will randomly generate **candidate samples** from some simple distribution, and then decide to accept or reject the candidate

Metropolis-Hastings sampling

- A widely applicable algorithm for numerical integration is the **Metropolis-Hastings** (MH) algorithm
- In the MH algorithm, we will randomly generate **candidate samples** from some simple distribution, and then decide to accept or reject the candidate
- MH algorithms need some customization and fine-tuning to be most efficient

Metropolis-Hastings sampling

- A widely applicable algorithm for numerical integration is the **Metropolis-Hastings** (MH) algorithm
- In the MH algorithm, we will randomly generate **candidate samples** from some simple distribution, and then decide to accept or reject the candidate
- MH algorithms need some customization and fine-tuning to be most efficient
- One intuition for the MH algorithm is that it is a rejection sampler with an adaptive envelope

Metropolis-Hastings sampler

Given a function $f(\theta) \propto p(\theta|D)$, and a symmetric **candidate generating distribution** $Q(a|b) = Q(b|a)$, a Metropolis-Hastings sampling algorithm proceeds as follows:

- 1 **Set** $i \leftarrow 1$ and **choose** $R \leftarrow 1000$
- 2 **Choose**, arbitrarily, $\theta^{(0)}$
- 3 **Draw** a randomly selected θ^c from $Q(\theta|\theta^{(i-1)})$
- 4 **Compute** the acceptance ratio $\alpha = \frac{f(\theta^c)}{f(\theta^{(i-1)})} = \frac{p(\theta^c|D)}{p(\theta^{(i-1)}|D)}$
- 5 **Draw** a randomly selected u from $U(0, 1)$. **If** $\alpha > u$, **set** $\theta^{(i)} \leftarrow \theta^c$, **otherwise set** $\theta^{(i)} \leftarrow \theta^{(i-1)}$
- 6 **Set** $i \leftarrow i + 1$. **If** $i \leq R$, **return** to Step 3, **otherwise halt**

Metropolis-Hastings sampler

- Because the accepted values may approach the correct distribution only slowly, we have to make sure the chain of samples has converged to a stationary sampling state before using the samples for inference

Metropolis-Hastings sampler

- Because the accepted values may approach the correct distribution only slowly, we have to make sure the chain of samples has converged to a stationary sampling state before using the samples for inference
- Specifically, we will discard a number of initial samples known as the **burn-in**:

$$\hat{\beta}_0 = \frac{1}{R - B} \sum_{i=B}^R \beta_0^{(i)}$$

Metropolis-Hastings sampler

- Because the accepted values may approach the correct distribution only slowly, we have to make sure the chain of samples has converged to a stationary sampling state before using the samples for inference
- Specifically, we will discard a number of initial samples known as the **burn-in**:

$$\hat{\beta}_0 = \frac{1}{R - B} \sum_{i=B}^R \beta_0^{(i)}$$

- We also repeat the procedure a few times with different values for $\theta^{(0)}$ to ensure that the algorithm converges to the same stationary distribution. Several convergence statistics exist, with Geweke's and Gelman's \hat{R} being the most popular

Sampling in MATLAB

Exercise: Metropolis sampling in MATLAB

Write a new class called `Metropolis()`

The constructor should work by taking an anonymous function as input (call the property `Target` of type `function_handle`), with other inputs such as parameters of the algorithm optional (with defaults)

In addition to the standard methods, `Metropolis()` should have at least a method `DrawSamples(N)` to draw N samples

You should assume that `TargetLogPdf` has exactly one input, but the input may be a vector

Finally, make sure the `disp()` method outputs something informative about the current state of the sampler